

# Kubernetes cluster network security

---

Jacopo Bufalino, Aalto University



Aalto University  
School of Science

12/09/2024

# Agenda

- **Kubernetes key concepts**
- **State of the art of kubernetes (network) security**
- **Cluster network misconfigurations**
  - Find
  - Collect and Analyze
  - Mitigate
- **Takeaways and lessons learned**

# Key concepts



kubernetes

Platform to run containerized workloads

Configuration as Code (YAML)

Used on-premises, on hybrid and cloud infrastructures

Abstracts concepts to favour portability (e.g., CRI)

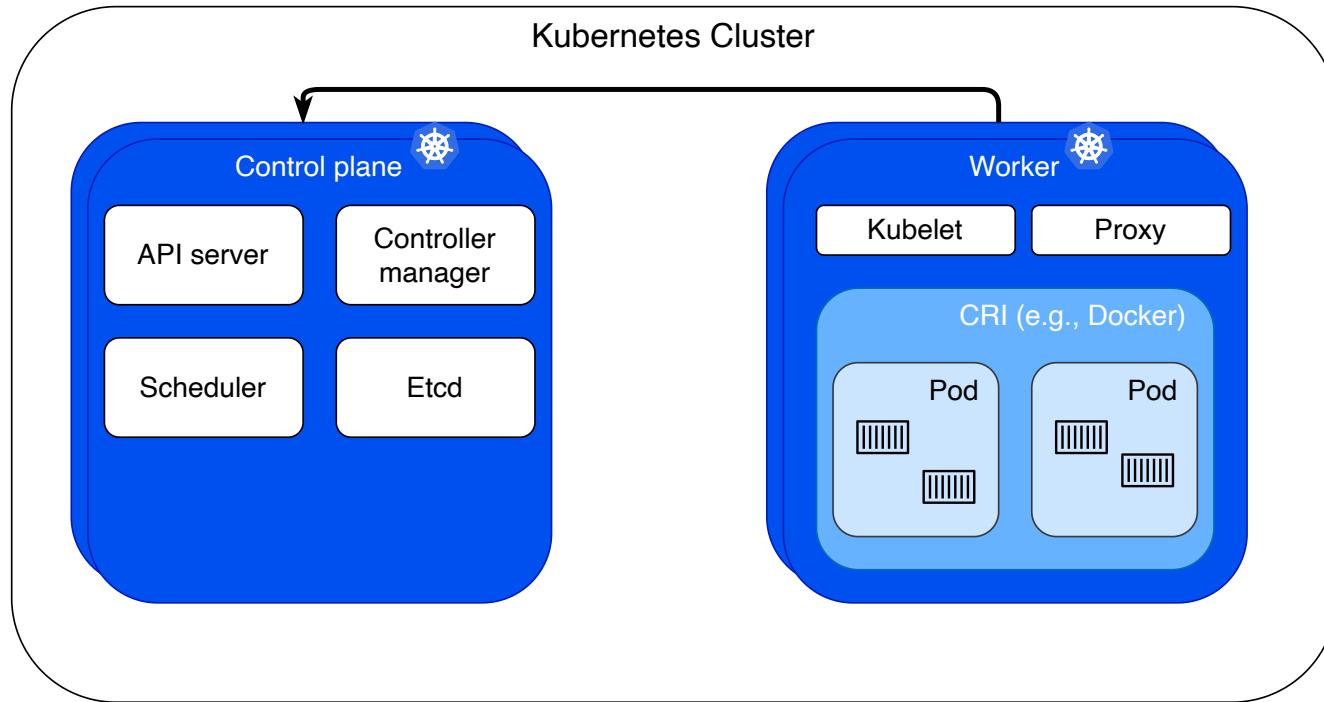


Template engine and package manager

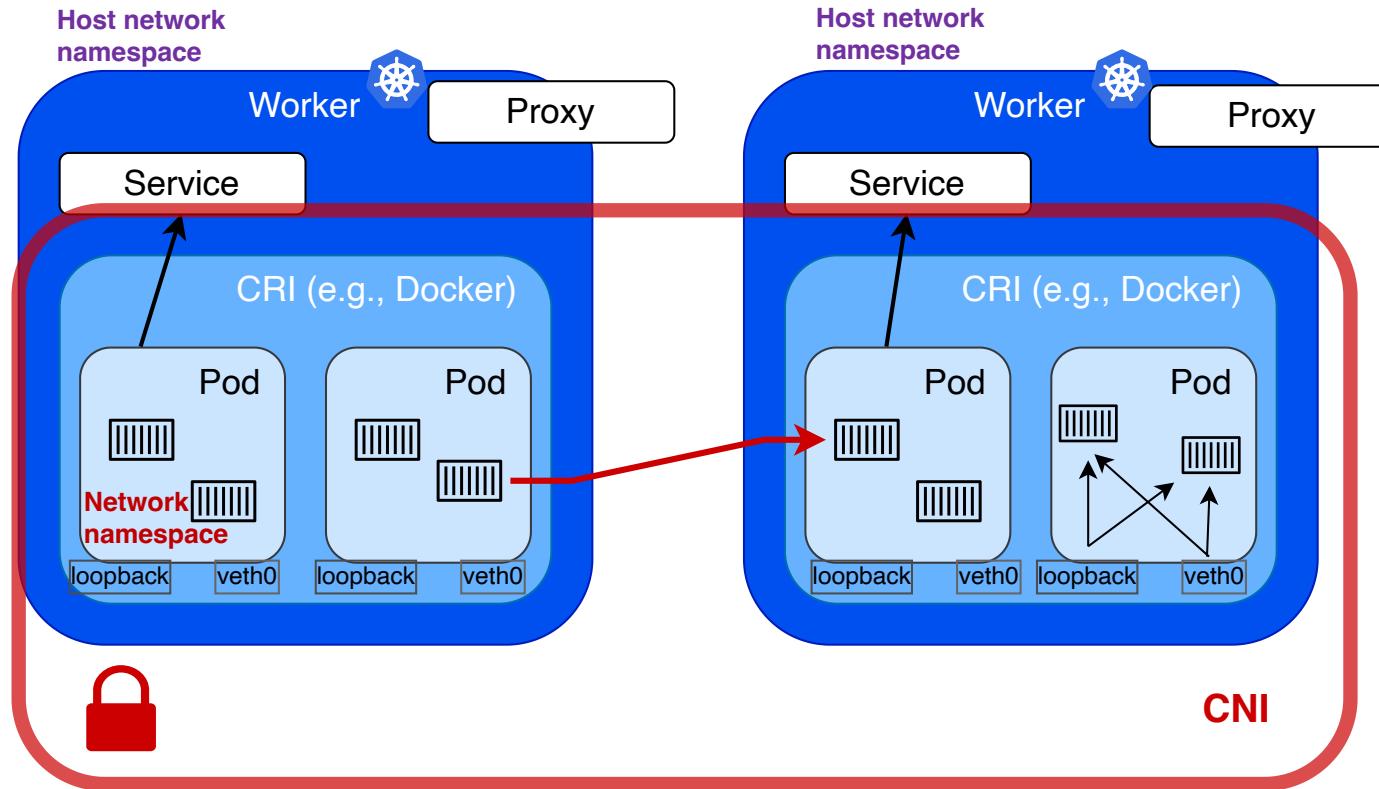
Share and create applications (charts)

Artifacthub.io has the largest collection of charts

# Kubernetes architecture (overview)



# Kubernetes networking model



+33-15669

Search the blog

Products Solutions Why Akamai Resources Partners Contact Us

Blog > Security Research >

Can't Be Contained: Finding a Command Injection Vulnerability in Kubernetes

## Can't Be Contained: Finding a Command Injection Vulnerability in Kubernetes



Research Cloud security Microsoft Defender for Cloud Cloud threats · 4 min read

Attackers exploiting new critical OpenMetadata vulnerabilities on Kubernetes clusters

By Microsoft Threat Intelligence

## Kubernetes Exposed: One Yaml away from Disaster

Threat Alert

Hackers like Kubernetes too...

Solutions Why Akamai Resources Partners

What a Cluster: Local Volumes Vulnerabilit

# MITRE® Threat matrix for Kubernetes

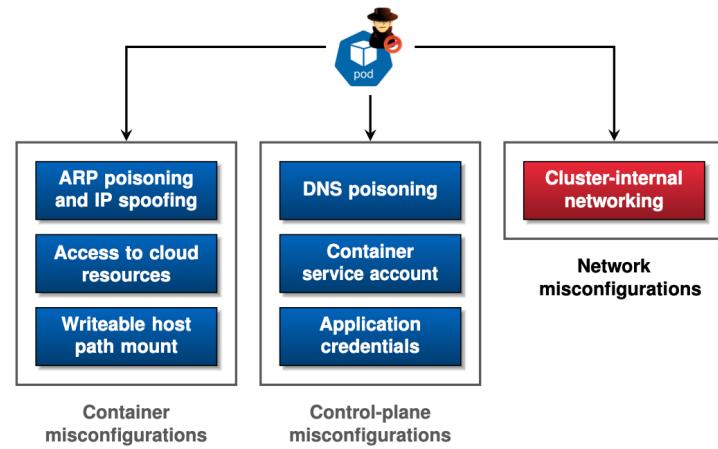
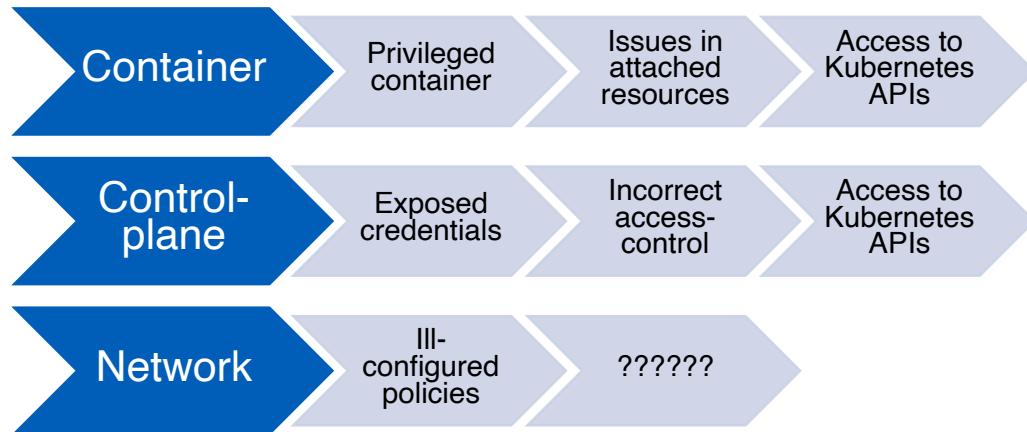
## Tactics

Tactics
Initial Access >
Execution >
Persistence >
Privilege Escalation >
Defense Evasion >
Credential Access >
Discovery >
Lateral Movement >
Collection >
Impact >

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access Kubernetes API server	Access cloud resources	Images from a private registry	Data destruction
Compromised image in registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Collecting data from pod	Resource hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from proxy server	Application credentials in configuration files	Exposed sensitive interfaces	Application credentials in configuration files		
Exposed sensitive interfaces	SSH server running inside container	Container service account			Access managed identity credentials	Instance Metadata API	Writable hostPath mount		
	Sidecar injection	Static pods			Malicious admission controller		CoreDNS poisoning		
							ARP poisoning and IP spoofing		

# MITRE® Threat matrix for Kubernetes

The 7 lateral movement techniques



MITRE® description:

[...]Attackers who gain access to a single container may use it for network reachability to another container in the cluster.

**HOW?  
STILL UNCLEAR !**

# Kubernetes cluster network security: state of the art

Static analysis of YAML files



Analysis using Kubernetes APIs



Kubescape



kube-hunter



Security platforms



StackRox

# Kubernetes cluster network security: state of the art

Misconfigurations reported

↳ Lack of network policies

↳ Container attached to **host network**

What is missing

↳ **Connections** between resources

↳ Report of **potential** connectivity

# Example of undetected misconfiguration

NGINX Pod is available to the cluster via a Service. A NetworkPolicy is in place to avoid unwanted access.

## Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - name: nginx-port
          containerPort: 80
          protocol: TCP
```

## Service

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    name: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: nginx-port
    - protocol: TCP
      port: 8081
      targetPort: 81
```

## NetworkPolicy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: nginx-policy
spec:
  podSelector:
    matchLabels:
      name: nginx
...
  ports:
    - protocol: TCP
      port: 6379
    - protocol: TCP
      port: nginx-port
```

# Goal #1: explore the internal cluster network

Network policies and services can be malformed

- How to ensure their correct application?

Observability

- Can we monitor the potential connectivity between applications?

Container ports declarations are only documentation

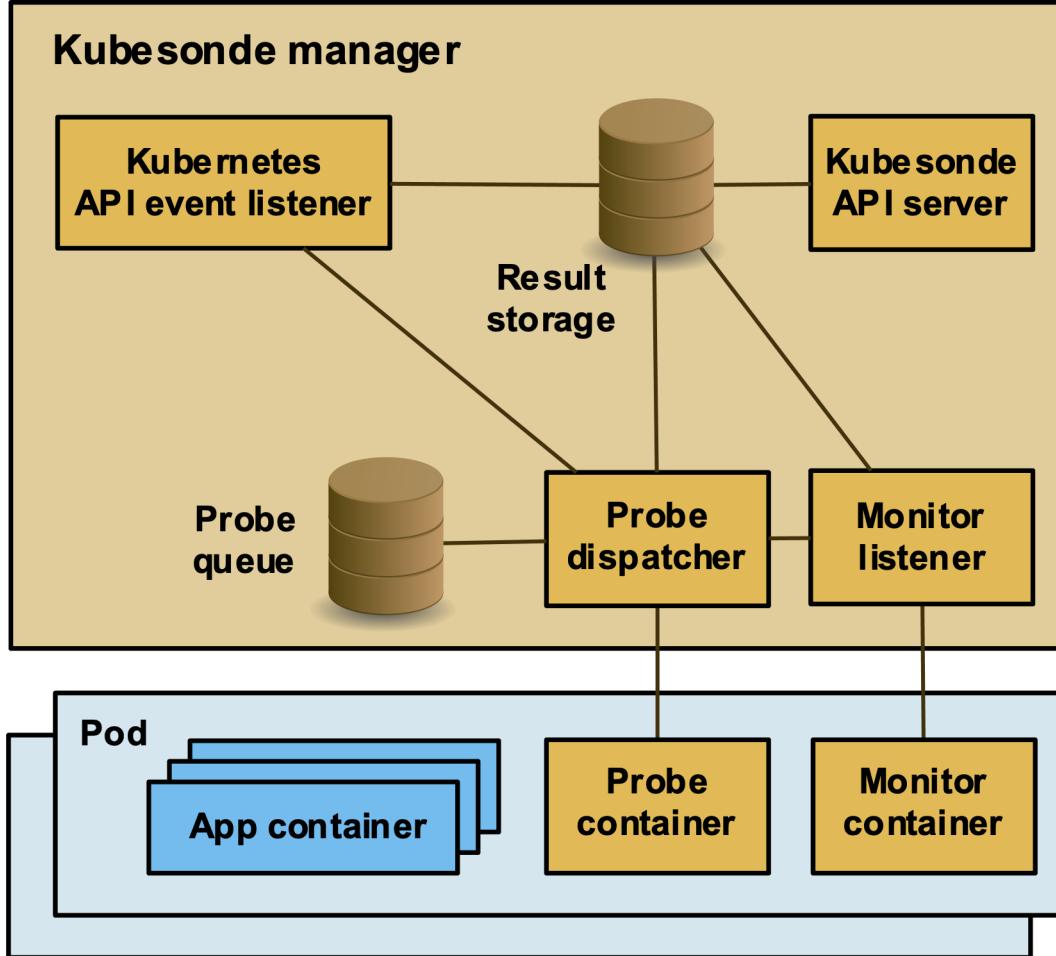
- Is that really the case?

# Kubesonde



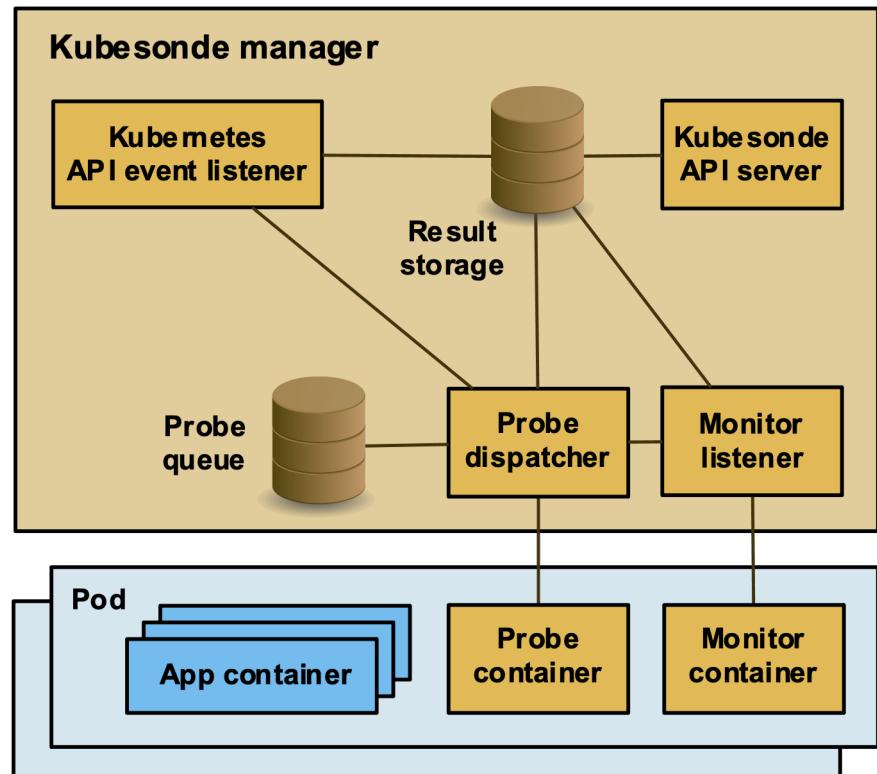
# Kubesonde

How it works



# Kubesonde

## How it works



# Testbed and results

Run Kubesonde on the 20 most popular Helm charts

Open ports not matching specifications

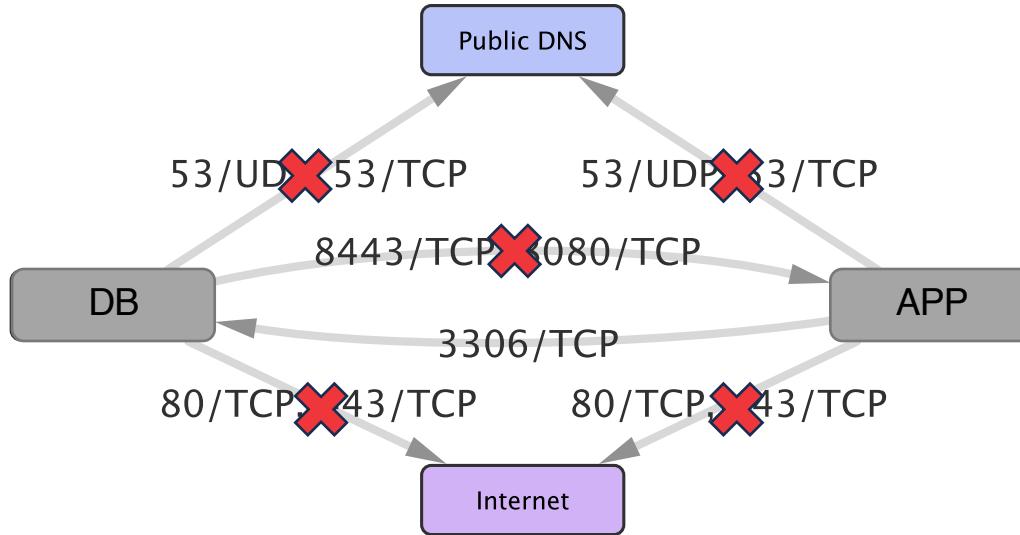
Containers listening on all interfaces

Lack of access control between pods

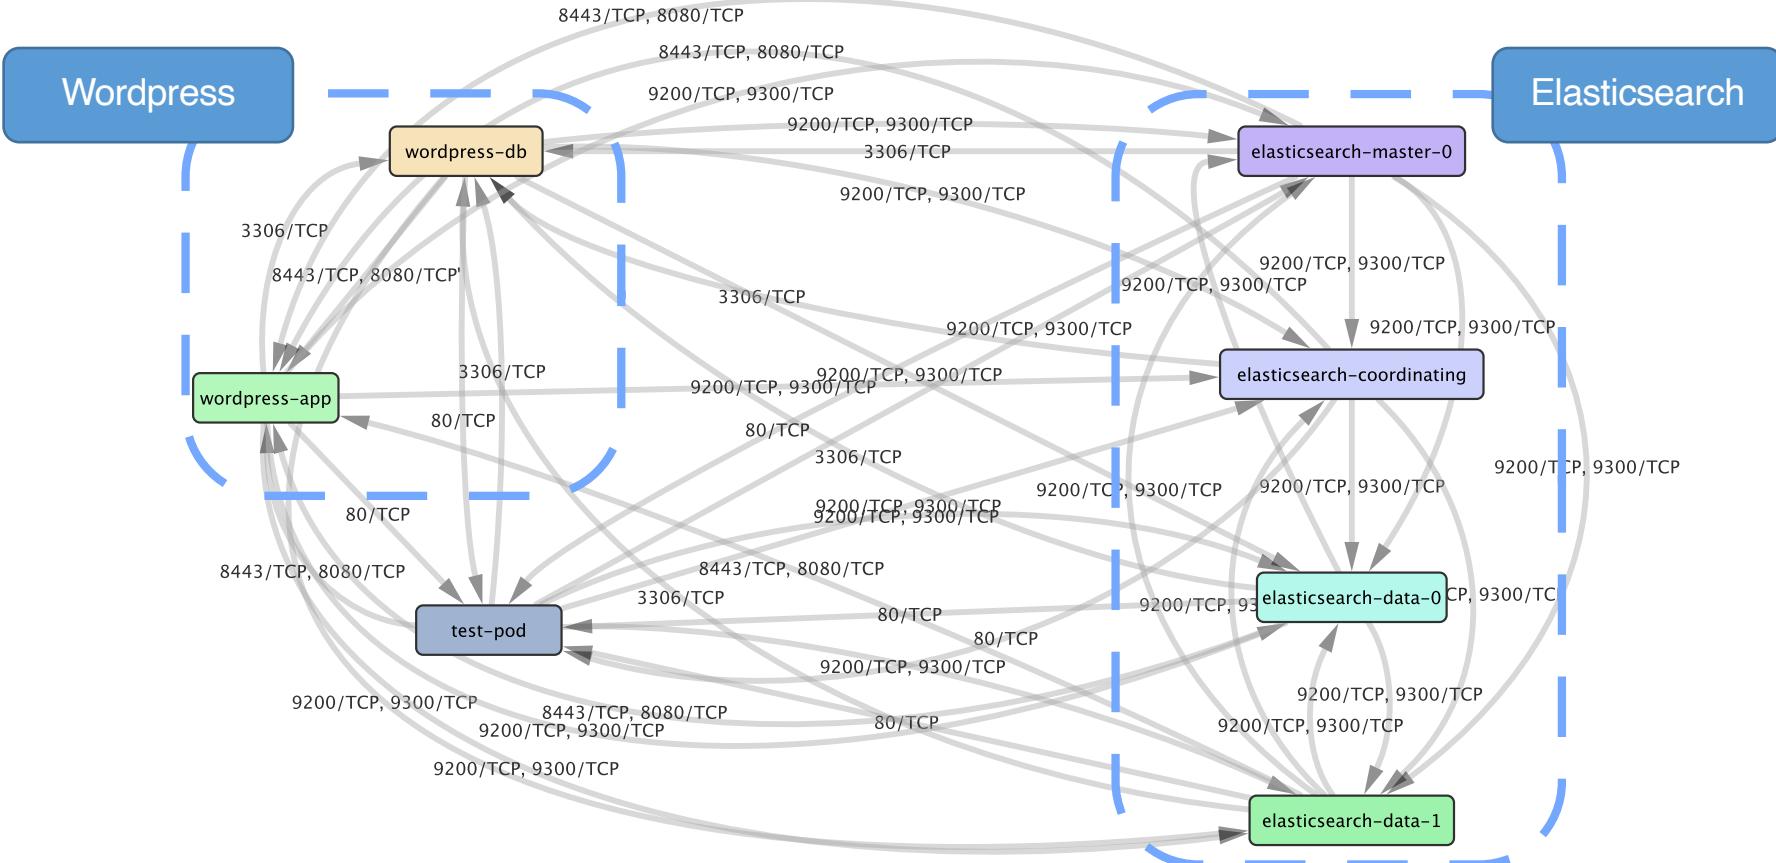
Access to public DNS and Internet

Dynamic ports

# Results - simple application



# Results - service composition



# Where we are so far...

Many types of misconfigurations

- More precise mapping is needed

Network policies are usually not available

- Developers must have knowledge of the applications

# Mapping misconfigurations

# Goal #2: map misconfigurations in the internal cluster network

Study their impact on cluster security

Investigate how frequently they arise

# The misconfigurations

M1

Port open on container is not declared

```
containers :  
- name : flink  
image : bitnami/flink  
ports :  
- containerPort : 8081
```

```
# netstat -a  
Active Internet connections ( servers and established )  
Proto Recv-Q Send-Q Local Address Foreign Address State  
tcp 0 0 0.0.0.8081 0.0.0.0:* LISTEN  
tcp 0 0 0.0.0.8080 0.0.0.0:* LISTEN
```

M2

Container allocates dynamic ports

```
containers :  
- name : clickhouse  
image :  
bitnami/clickhouse  
ports :  
- containerPort : 8080
```

```
# netstat -a  
Active Internet connections ( servers and established )  
Proto Recv-Q Send-Q Local Address Foreign Address State  
tcp 0 0 0.0.0.8080 0.0.0.0:* LISTEN  
tcp 0 0 0.0.0.47123 0.0.0.0:* LISTEN
```

# The misconfigurations

M3

Port declared on container is not open

containers :

- name : consul

image : bitnami/consul

ports :

- containerPort : 8080

- containerPort : 8300

- containerPort : 8360

# netstat -a

Active Internet connections ( servers and established )

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State
tcp	0	0	0.0.0.0:8080	0.0.0.0:*	LISTEN

M4

Label collision

# Pod 1

...

labels:  
role: frontend

...

# Pod 2

...

labels:  
role: frontend

...

# The misconfigurations

M5

Service issues

M6

Lack of network policies

M7

Container binds to host network

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    name: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: nginx-port
    - protocol: TCP
      port: 8081
      targetPort: 81
```

M5

```
apiVersion: v1
kind: Pod
metadata:
  name: influxdb
spec:
  hostNetwork: true
  containers:
    - name: influxdb
      image: influxdb
```

M7

# The data



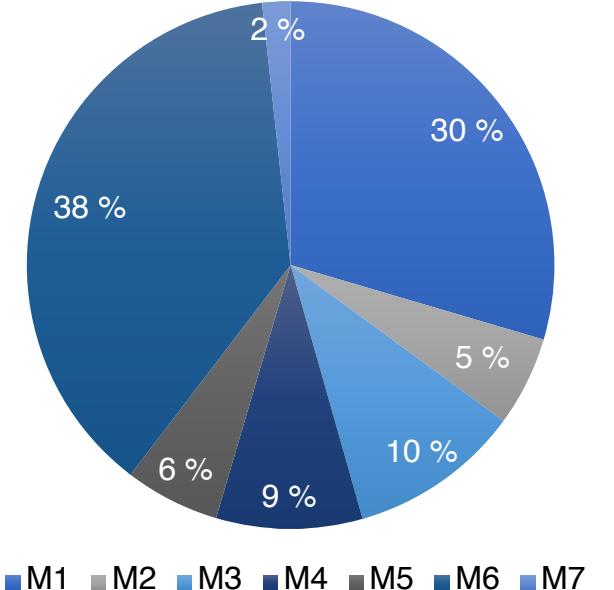
- Selected **6** datasets from Helm providers including:
  - **Bitnami**
  - **Wikimedia**
  - **CNCF**
  - **Prometheus Community**
  - **European Environmental Agency**
  - **Banzaicloud**
- **287** charts in total

# Findings

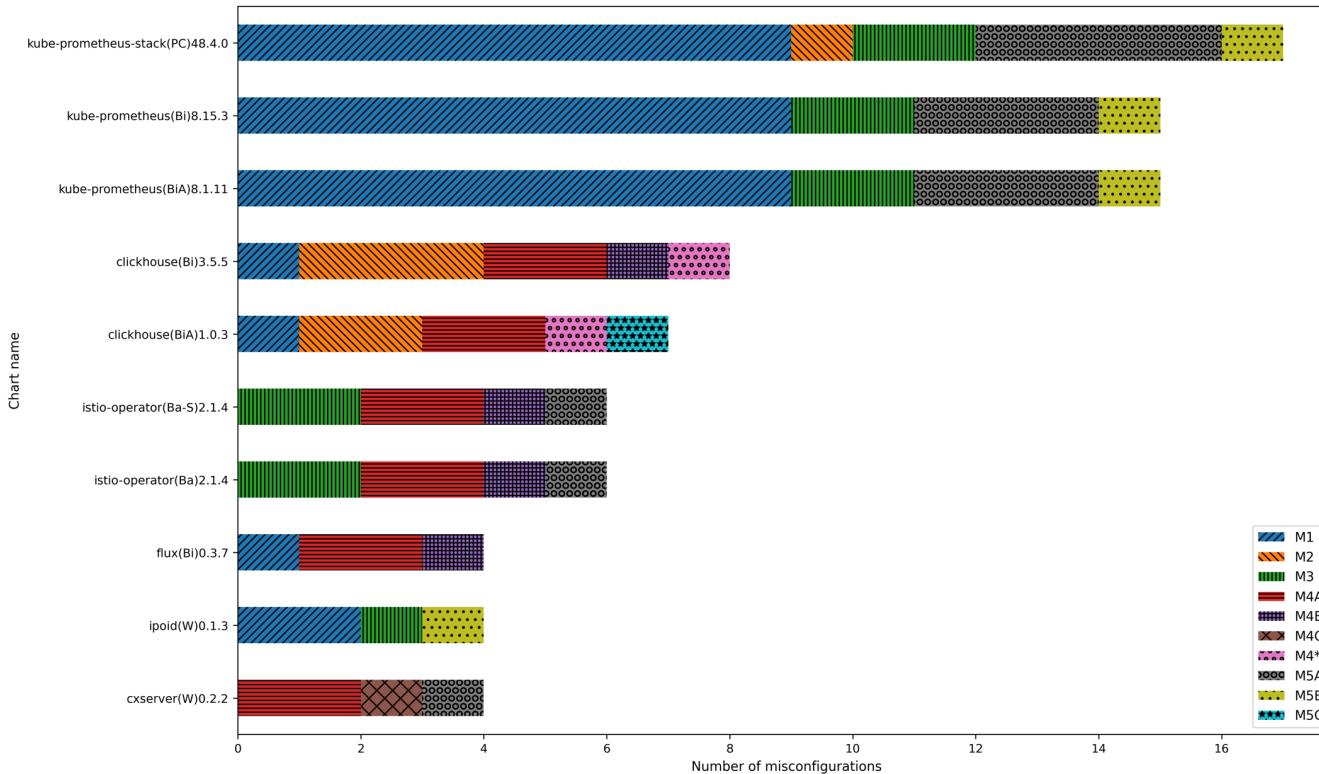
Over 90% of the charts have at least one misconfiguration (259 / 287)

**Misconfigurations (634 in total)**

- More than 80% (241/287) of the apps do not include NetworkPolicy
- 30% of the total are type M1 (Port not declared)
- 60% not reported by state-of-the-art tools
- 8% of the charts with NetworkPolicy have errors



# The top 10



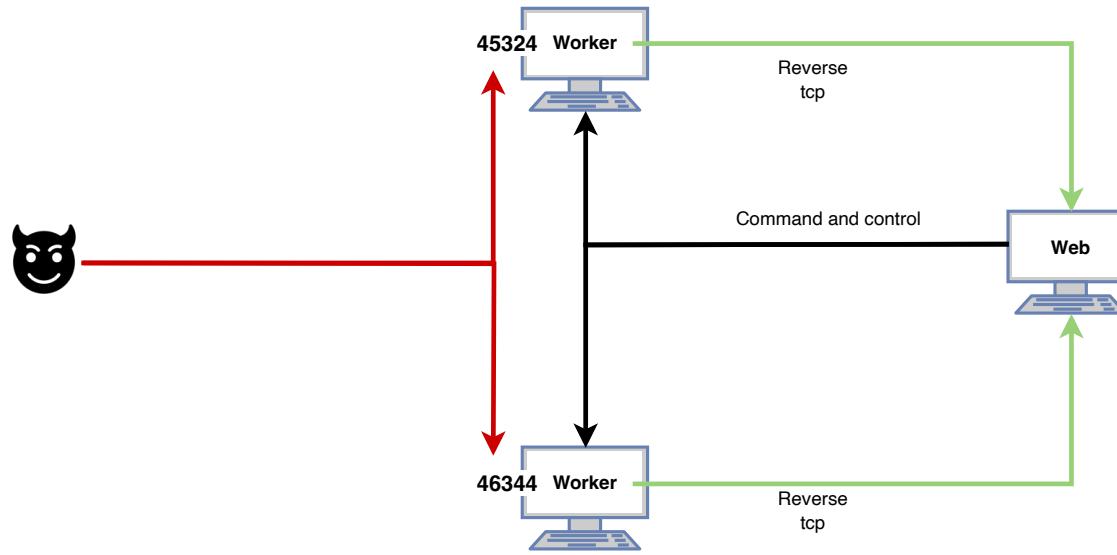
# Successful attack



Concourse

Concourse is a CI/CD software part of the CNCF landscape

Master web node and a number of different workers



[bitnami/flux] fix: 🐛 🔒 Add missing notification controller ports #25

\_merged javsalgar merged 2 commits into main from fix/flux-add-missing-ports 2 weeks ago

 Conversation 7     Commits 2     Checks 11     Files changed 7

 **javasalgar** commented 3 weeks ago Contributor • ...

[bitnami/jaeger] fix: 🐛 🔒 Expose missing ports in deployment spec

Merged javsalgar merged 1 commit into main from fix/jaeger-add-missing-ports last week

---

Conversation 0    Commit 1    Checks 20    Files changed 8

**Many organizations fix misconfiguration**

javasalgar commented 2 weeks ago

Signed-off-by: Javier Salmeron Garcia [js](#)

Description of the change

This PR adds to the deployment spec the values.

 Wikimedia Code Review    CHANGES ▾    DOCUMENTATION ▾    BROWSE ▾

```
26    {{- end }}
27    serviceAccountName: {{ include "app.serviceAccountName" -}}
28
29    containers:
30      - name: {{ .Release.Name }}
31      image: {{( .Values.image.repository )|(.Values.image.tag | default .Chart.AppVersion) }}
32
33    ports:
34      - name: nessus
35        containerPort: 8884
36        protocol: TCP
37
38    # Port 8884 - User Interface, TSC communication, and API calls
39
40    ports:
41      - name: nessus
42        containerPort: 8884
43        protocol: TCP
44
45    stdIn: true
46    stdOut: true
47
48    # Port 8884 - User Interface, TSC communication, and API calls
49
50    ports:
51      - name: nessus
52        containerPort: 8884
53        protocol: TCP
54
55    {{- with .Values.nodeSelector -}}
56    nodeSelector:
57      {{- toyaml . | indent 8 -}}
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131    - name: serfwan-udp
132      containerPort: {{ .Values.containerPorts.serfLAN }}
133      protocol: "UDP"
134    + - name: serfwan-tcp
135      containerPort: {{ .Values.containerPorts.serfLAN }}
136      protocol: "TCP"
137    + - name: serfwan-udp
138      containerPort: {{ .Values.containerPorts.serfWAN }}
139      protocol: "UDP"
140    + - name: rpc-server
141      containerPort: {{ .Values.containerPorts.rpcServer }}
142      - name: dns-tcp
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
```

- bitnami/charts #25106 [bitnami/kafka] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25141 [bitnami/zookeeper] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25069 [bitnami/contour] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25078 [bitnami/grafana-operator] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25156 [bitnami/jaeger] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25097 [bitnami/grafana-tempo] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25110 [bitnami/keycloak] fix: 🐛 🔒 Expose missing ports in deployment spec and fix headless service
- bitnami/charts #25113 [bitnami/kube-state-metrics] fix: 🐛 🔒 Expose missing ports in deployment spec

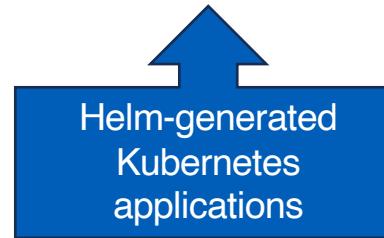
- [bitnami/mysql] fix: 🐛 Add support for mysqlx port
- bitnami/charts #25077 [bitnami/grafana-loki] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25043 [bitnami/argo-cd] fix: 🐛 🔒 Expose metrics port in deployment definition
- bitnami/charts #25042 [bitnami/appsmith] fix: 🐛 🔒 Add ambassador container to appsmith-backend to contact appsmith-rts
- bitnami/charts #25072 [bitnami/ejbcal] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25066 [bitnami/consul] fix: 🐛 🔒 Expose missing ports in deployment spec
- bitnami/charts #25048 [bitnami/clickhouse] fix: 🐛 🔒 Add shard label to avoid Compute Unit collision
- bitnami/charts #25045 [bitnami/cassandra] fix: 🐛 🔒 Do not expose tls internode port unless encryption is set
- bitnami/charts #25041 [bitnami/apisix] fix: 🐛 🔒 Do not expose http-metrics unless metrics.enabled=true
- bitnami/charts #25047 [bitnami/cert-manager] fix: 🐛 🔒 Expose missing ports in deployment spec



Aalto University  
School of Science

# Mitigate misconfigurations (HelmET)

# Goal #3: create secure network policies for unknown charts



## Challenges

Not all applications are open-source

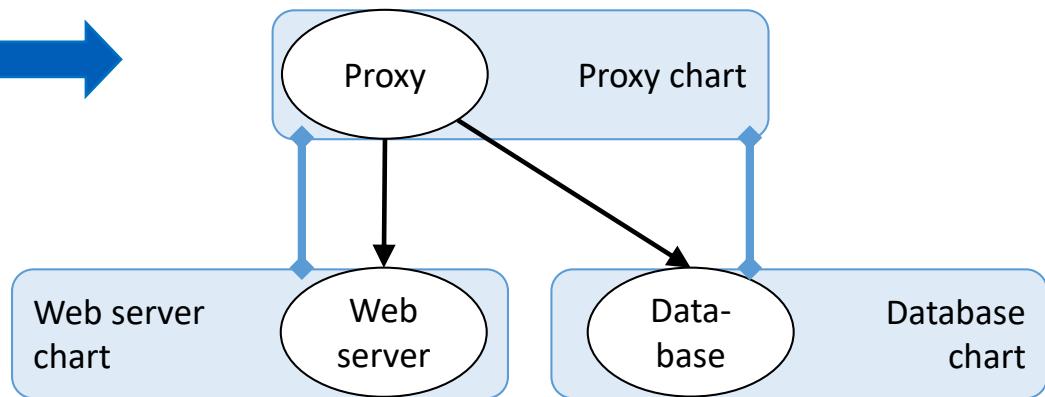
How to create policies before deployment?

# Idea

Generate policies using the chart structure and dependencies



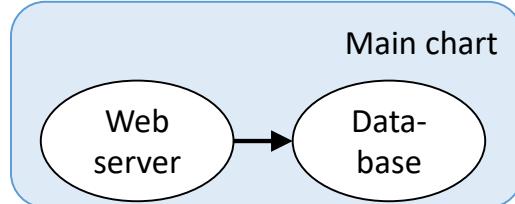
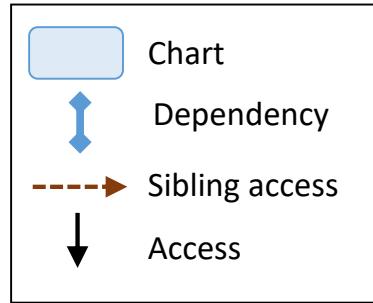
YAML



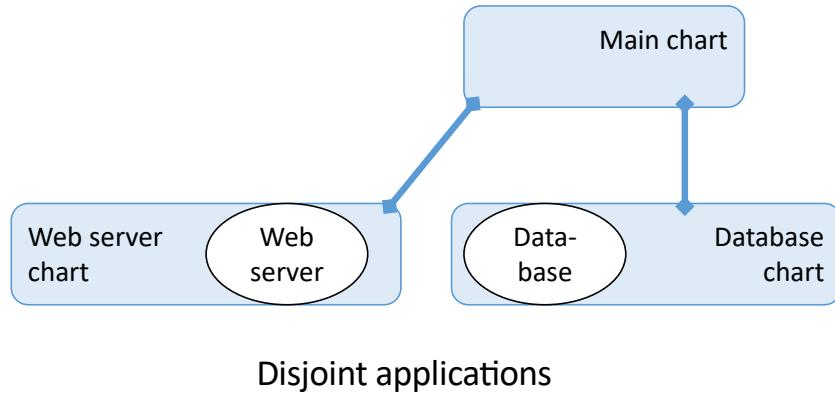
# Dataset and rule generation

- We analyzed **465** charts
- Identified **5** composition patterns
- Generate policies using software engineering concepts
  - Information hiding
  - Modularity
  - Module hierarchy

# Dataset and rule generation

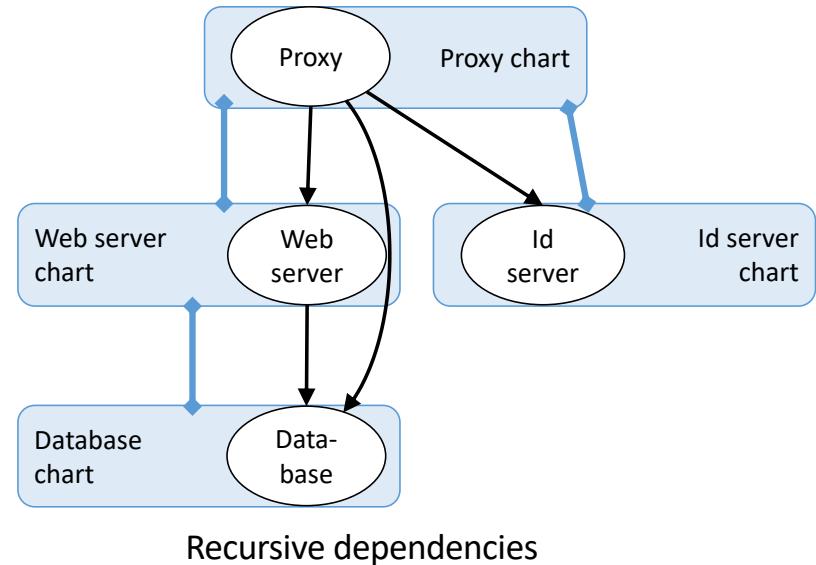
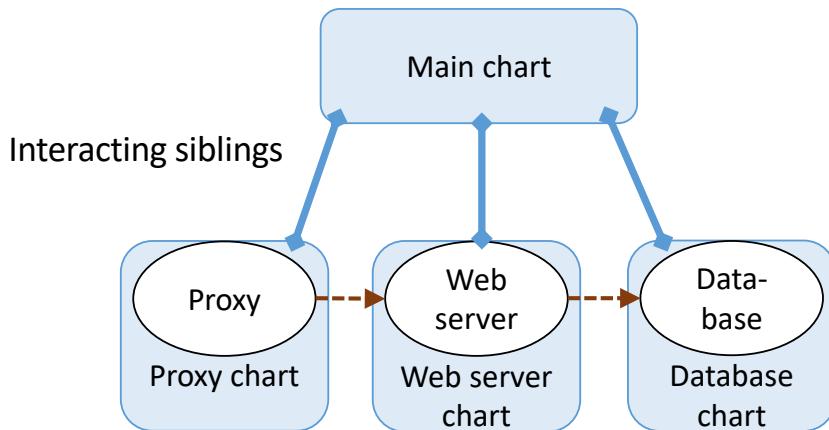
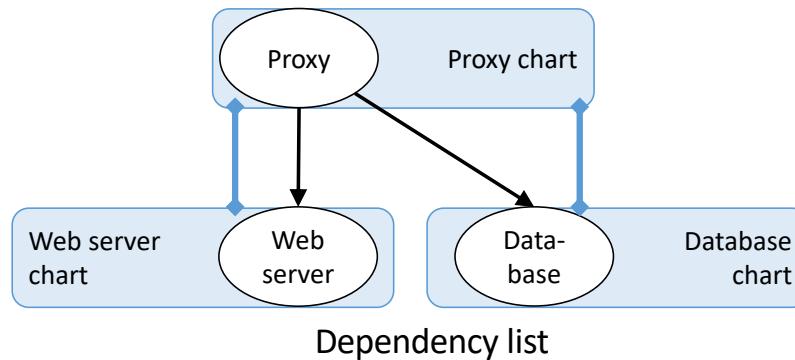


Standalone application



Disjoint applications

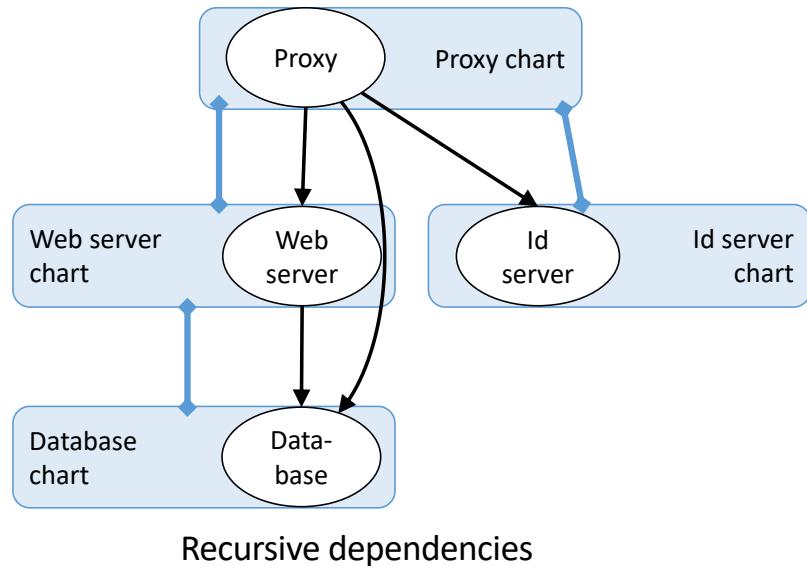
# Dataset and rule generation



# Dataset and rule generation

Three simple rules:

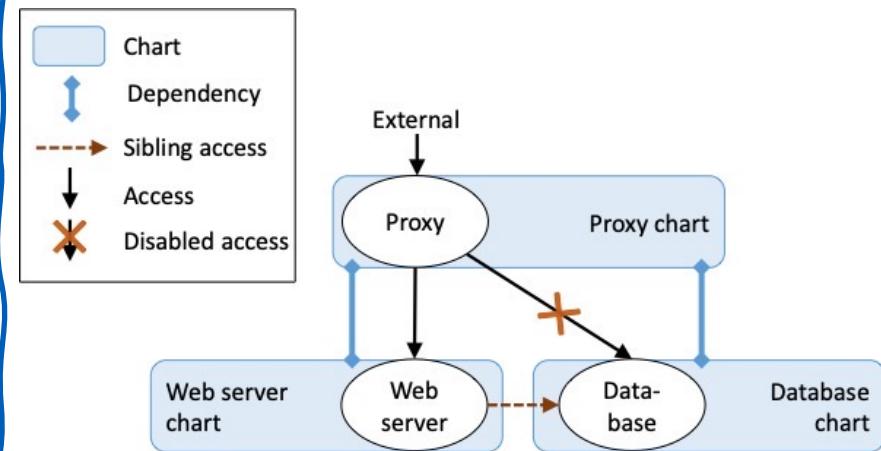
1. Deny access between siblings
2. Allow access only from parent to children
3. Default-deny all other connections



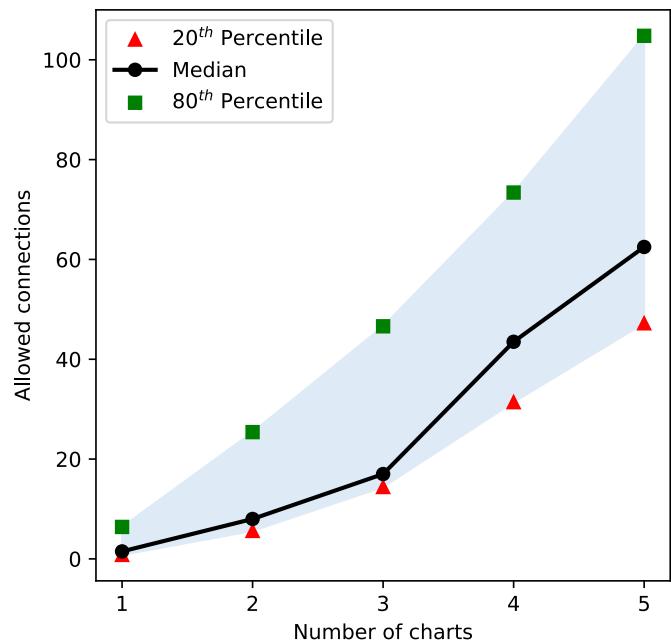
# Results

- Over 90% of the applications can be automatically secured
- **High level policy language** for the remaining ones

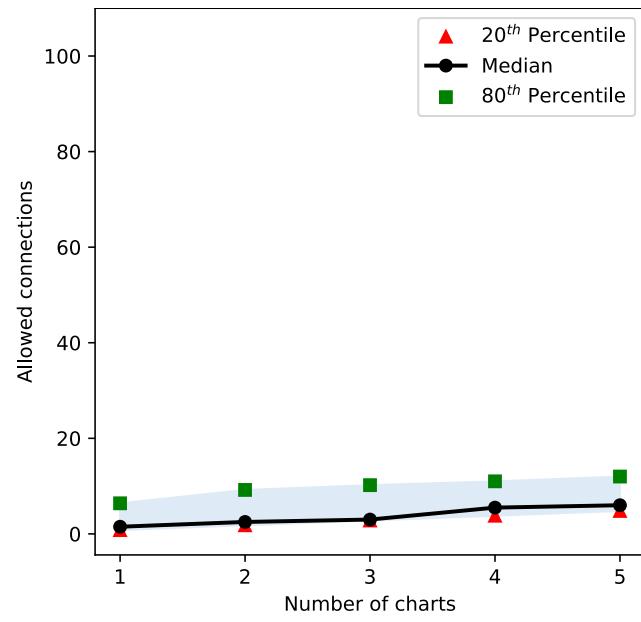
```
Name: "proxy-policy"
ComponentSelector:
...
Interactions:
  - From:
    chart: "proxy"
  To:
    chart: "webserver"
# Enable communication between siblings:
  - From:
    chart: "webserver"
  To:
    chart: "database"
# Disable direct access from proxy to database:
# - From:
#   chart: "proxy"
# To:
#   chart: "database"
...
...
```



# Results

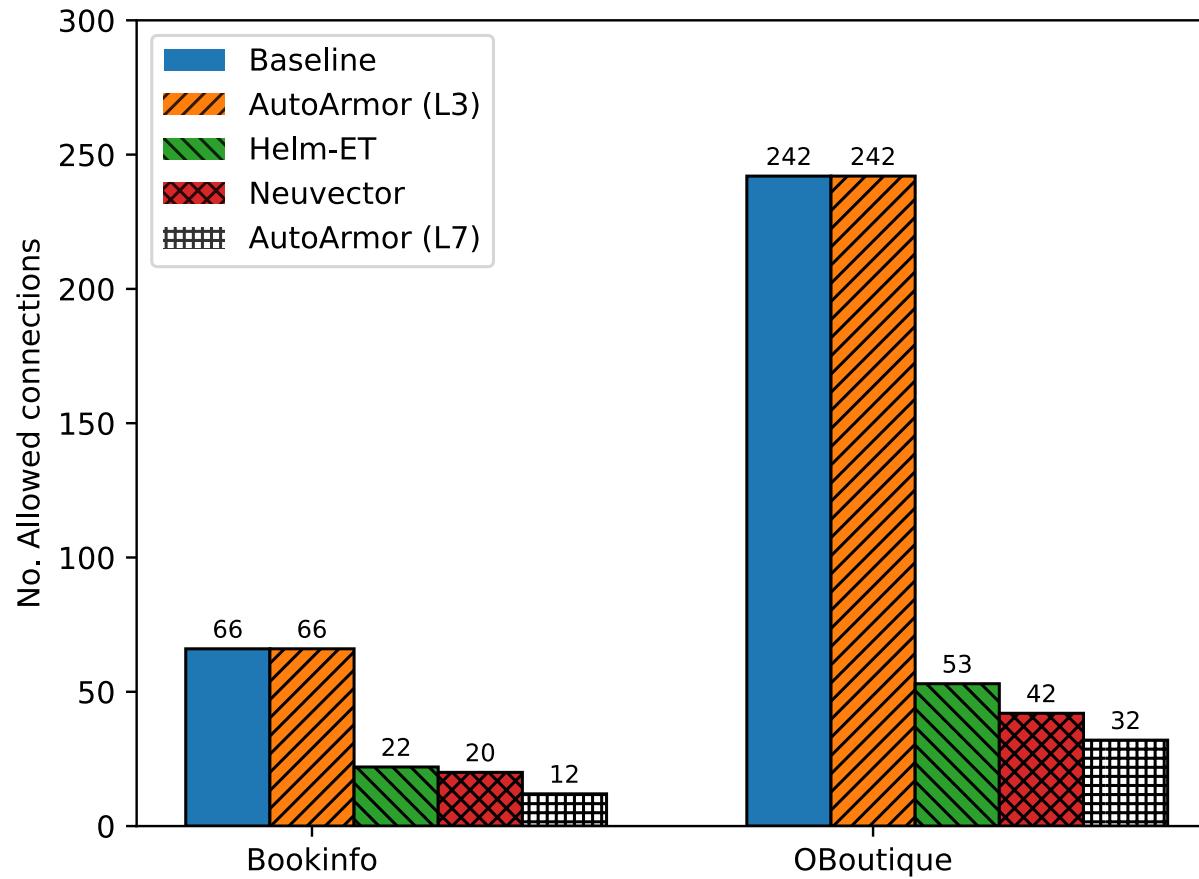


Before



After

# Comparison with other work



# Takeaways

- Still long way for lateral movement protection/mitigation
- Do not assume third party applications to be hardened in terms of internal networking
- Much depends on the developers to ensure proper access control

# The end

Jacopo Bufalino



@JacopoBufalino



github.com/jackap



<https://github.com/jackap/kubesonde>



Aalto University  
School of Science

Cluster network  
Security survey

